

Investigate_a_Dataset

October 20, 2018

Tip: Welcome to the Investigate a Dataset project! You will find tips in quoted sections like this to help organize your approach to your investigation. Before submitting your project, it will be a good idea to go back through your report and remove these sections to make the presentation of your work as tidy as possible. First things first, you might want to double-click this Markdown cell and change the title so that it reflects your dataset and investigation.

1 Project: Investigate a Dataset (TMDb_Movies Dataset)

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

For third Data Analysis project: I have selected TMDb movies dataset. This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue. It consists of 21 columns such as imdb_id, revenue, budget, vote_count etc.

The questions that I am interested to analyze from this dataset: * What the successful movie genres? * What the relationship between the popularity and the runtime of the movie? * What the relationship between the budget and the runtime of the movie? * what's the highest budget and profit in the TMDb_Movies dataset? * What's the lowest budget and profit in the TMDb_Movies dataset?

First: In this section of the report, Import the necessary package and use `pd.read_csv` to load the movie dataset, then print the first rows.

```
In [26]: # Use this cell to set up import statements for all of the packages that you
        #      plan to use.

        # Remember to include a 'magic word' so that your visualizations are plotted
        #      inline with the notebook. See this page for more:
        #      http://ipython.readthedocs.io/en/stable/interactive/magics.html
import numpy as np
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
% matplotlib inline
```

Data Wrangling

First: In this section of the report, Import the necessary package and use `pd.read_csv` to load the movie dataset, then print the first rows.

1.1.1 General Properties

```
In [27]: # Load your data and print out a few lines. Perform operations to inspect data
#        types and look for instances of missing or possibly errant data.
df= pd.read_csv('tmdb-movies.csv')
df.head()
```

```
Out[27]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	

3	Every generation has a story.	...
4	Vengeance Hits Home	...

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	
2	Beatrice Prior must confront her inner demons ...	119	
3	Thirty years after defeating the Galactic Empi...	136	
4	Deckard Shaw seeks revenge against Dominic Tor...	137	

	genres	\
0	Action Adventure Science Fiction Thriller	
1	Action Adventure Science Fiction Thriller	
2	Adventure Science Fiction Thriller	
3	Action Adventure Science Fiction Fantasy	
4	Action Crime Thriller	

	production_companies	release_date	vote_count	\
0	Universal Studios Amblin Entertainment Legenda...	6/9/15	5562	
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15	6185	
2	Summit Entertainment Mandeville Films Red Wago...	3/18/15	2480	
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15	5292	
4	Universal Pictures Original Film Media Rights ...	4/1/15	2947	

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

In [28]: df.shape

Out[28]: (10866, 21)

It has 10866 columns and 21 rows.

In [29]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
```

```

original_title      10866 non-null object
cast                10790 non-null object
homepage            2936 non-null object
director            10822 non-null object
tagline             8042 non-null object
keywords            9373 non-null object
overview            10862 non-null object
runtime             10866 non-null int64
genres              10843 non-null object
production_companies 9836 non-null object
release_date        10866 non-null object
vote_count          10866 non-null int64
vote_average        10866 non-null float64
release_year        10866 non-null int64
budget_adj          10866 non-null float64
revenue_adj         10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

```
In [30]: df.describe()
```

```

Out[30]:

```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

From The Movie Database (TMDb): We can see that: Certain columns, like ‘cast’ and ‘genres’, * contain multiple values separated by pipe (|) characters. There are some odd characters in the ‘cast’ column. Don’t worry about cleaning them. You can leave them as is. * The final two columns ending with “_adj” show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

1.1.2 Data Cleaning (Removing the unused columns from The Movie Database (TMDb))

In this step: I will do * Removing the duplicate record(if any). * Removing the rows with null values (if any). * Removing the unused columns from The Movie Database. * Removing the zeros from budget and the revenue columns.

First: I will remove the unused columns which are id, imdb_id,original_title, homepage, tagline, overview, production_companies,vote_count,vote_average, budget_adj, and revenue_adj.

```
In [31]: #After discussing the structure of the data and any problems that need to be
#cleaned, perform those cleaning steps in the second part of this section.
df.drop(['id', 'imdb_id','original_title', 'homepage', 'tagline', 'overview', 'producti
df.head()
```

```
Out[31]:
```

	popularity	budget	revenue	\
0	32.985763	150000000	1513528810	
1	28.419936	150000000	378436354	
2	13.112507	110000000	295238201	
3	11.173104	200000000	2068178225	
4	9.335014	190000000	1506249360	

	cast	director	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	George Miller	
2	Shailene Woodley Theo James Kate Winslet Ansel...	Robert Schwentke	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	J.J. Abrams	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	James Wan	

	keywords	runtime	\
0	monster dna tyrannosaurus rex velociraptor island	124	
1	future chase post-apocalyptic dystopia australia	120	
2	based on novel revolution dystopia sequel dyst...	119	
3	android spaceship jedi space opera 3d	136	
4	car race speed revenge suspense car	137	

	genres	release_date	release_year
0	Action Adventure Science Fiction Thriller	6/9/15	2015
1	Action Adventure Science Fiction Thriller	5/13/15	2015
2	Adventure Science Fiction Thriller	3/18/15	2015
3	Action Adventure Science Fiction Fantasy	12/15/15	2015
4	Action Crime Thriller	4/1/15	2015

Removing the duplicate record.

```
In [32]: # I have removed the duplicate record.
df.drop_duplicates(inplace=True)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 10 columns):
popularity      10865 non-null float64
budget          10865 non-null int64
revenue         10865 non-null int64
cast            10789 non-null object
director        10821 non-null object
keywords        9372 non-null object
runtime         10865 non-null int64
genres          10842 non-null object
release_date    10865 non-null object
release_year    10865 non-null int64
dtypes: float64(1), int64(4), object(5)
memory usage: 933.7+ KB

```

Removing the rows with null values

```

In [33]: # I have removed rows with null values
         df.dropna(inplace=True)
         df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 9306 entries, 0 to 10865
Data columns (total 10 columns):
popularity      9306 non-null float64
budget          9306 non-null int64
revenue         9306 non-null int64
cast            9306 non-null object
director        9306 non-null object
keywords        9306 non-null object
runtime         9306 non-null int64
genres          9306 non-null object
release_date    9306 non-null object
release_year    9306 non-null int64
dtypes: float64(1), int64(4), object(5)
memory usage: 799.7+ KB

```

Removing zeros from budget and revenue columns.

```

In [34]: value=['budget', 'revenue']

         df[value] = df[value].replace(0, np.NaN)

         df.dropna(subset = value, inplace = True)

```

Create a column called "profit"

```
In [35]: df.loc[:, 'Profit'] = (df['revenue'] - df['budget'])
df.head()
```

```
Out[35]:
```

	popularity	budget	revenue	\
0	32.985763	150000000.0	1.513529e+09	
1	28.419936	150000000.0	3.784364e+08	
2	13.112507	110000000.0	2.952382e+08	
3	11.173104	200000000.0	2.068178e+09	
4	9.335014	190000000.0	1.506249e+09	

	cast	director	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	George Miller	
2	Shailene Woodley Theo James Kate Winslet Ansel...	Robert Schwentke	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	J.J. Abrams	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	James Wan	

	keywords	runtime	\
0	monster dna tyrannosaurus rex velociraptor island	124	
1	future chase post-apocalyptic dystopia australia	120	
2	based on novel revolution dystopia sequel dyst...	119	
3	android spaceship jedi space opera 3d	136	
4	car race speed revenge suspense car	137	

	genres	release_date	release_year	\
0	Action Adventure Science Fiction Thriller	6/9/15	2015	
1	Action Adventure Science Fiction Thriller	5/13/15	2015	
2	Adventure Science Fiction Thriller	3/18/15	2015	
3	Action Adventure Science Fiction Fantasy	12/15/15	2015	
4	Action Crime Thriller	4/1/15	2015	

	Profit
0	1.363529e+09
1	2.284364e+08
2	1.852382e+08
3	1.868178e+09
4	1.316249e+09

Create function to return the highest value

```
In [36]: def highest(column):
    high= df[column].idxmax()
    value=pd.DataFrame(df.loc[high])

    return value
```

```
In [37]: highest('budget')
```

```
Out[37]:
```

	2244
popularity	0.25054
budget	4.25e+08
revenue	1.10876e+07
cast	Kate Bosworth Jang Dong-gun Geoffrey Rush Dann...
director	Sngmoo Lee
keywords	assassin small town revenge deception super speed
runtime	100
genres	Adventure Fantasy Action Western Thriller
release_date	12/2/10
release_year	2010
Profit	-4.13912e+08

```
In [38]: highest('Profit')
```

```
Out[38]:
```

	1386
popularity	9.43277
budget	2.37e+08
revenue	2.78151e+09
cast	Sam Worthington Zoe Saldana Sigourney Weaver S...
director	James Cameron
keywords	culture clash future space war space colony so...
runtime	162
genres	Action Adventure Fantasy Science Fiction
release_date	12/10/09
release_year	2009
Profit	2.54451e+09

The highest value in the budget of movies was 4.25e+08. The highest value in the profit of movies was 2.54451e+09.

Create function to return the lowest value

```
In [39]: def lowest(column):
        low= df[column].idxmin()
        value=pd.DataFrame(df.loc[low])

        return value
```

```
In [40]: lowest('budget')
```

```
Out[40]:
```

	2618
popularity	0.090186
budget	1
revenue	100
cast	David Spade Sophie Marceau Ever Carradine Step...
director	Jeff Pollack


keywords	restaurant neighbor dog ring
runtime	95
genres	Comedy Romance
release_date	4/23/99
release_year	1999
Profit	99

In [41]: lowest('Profit')

Out[41]:	2244
popularity	0.25054
budget	4.25e+08
revenue	1.10876e+07
cast	Kate Bosworth Jang Dong-gun Geoffrey Rush Dann...
director	Sngmoo Lee
keywords	assassin small town revenge deception super speed
runtime	100
genres	Adventure Fantasy Action Western Thriller
release_date	12/2/10
release_year	2010
Profit	-4.13912e+08

The lowest value in budget was 1. The lowest value in profit was -4.13912e+08.

In [42]: data = df.groupby('release_year').mean()
data.hist(color='#86bf91',grid=False,);



Here, I just want to know the mean of the values of the release year.
Exploratory Data Analysis

Before answering the first question: let's create a function to split the values separated by pipe (|) characters.

```
In [43]: def count(column):
          count = df[column].str.cat(sep = '|')

          count = pd.Series(count.split('|'))

          x = count.value_counts(ascending = False)

          return x
```

1.1.3 Research Question 1 (Replace this header name!)


- What the successful movie genres?

```
In [44]: # Use this, and more code cells, to explore your data. Don't forget to add
          #   Markdown cells to document your observations and findings.
          x = count('genres')
          x.head()
```

```
Out[44]: Drama      1686
          Comedy     1301
          Thriller   1164
          Action     1050
          Adventure   729
          dtype: int64
```

```
In [52]: #Create bars and choose color
          #Add title and axis names
          x.plot(kind = "bar", color='pink')
          plt.title('The successful movie genres')
          plt.xlabel('categories')
          plt.ylabel('Number of Moives')
```

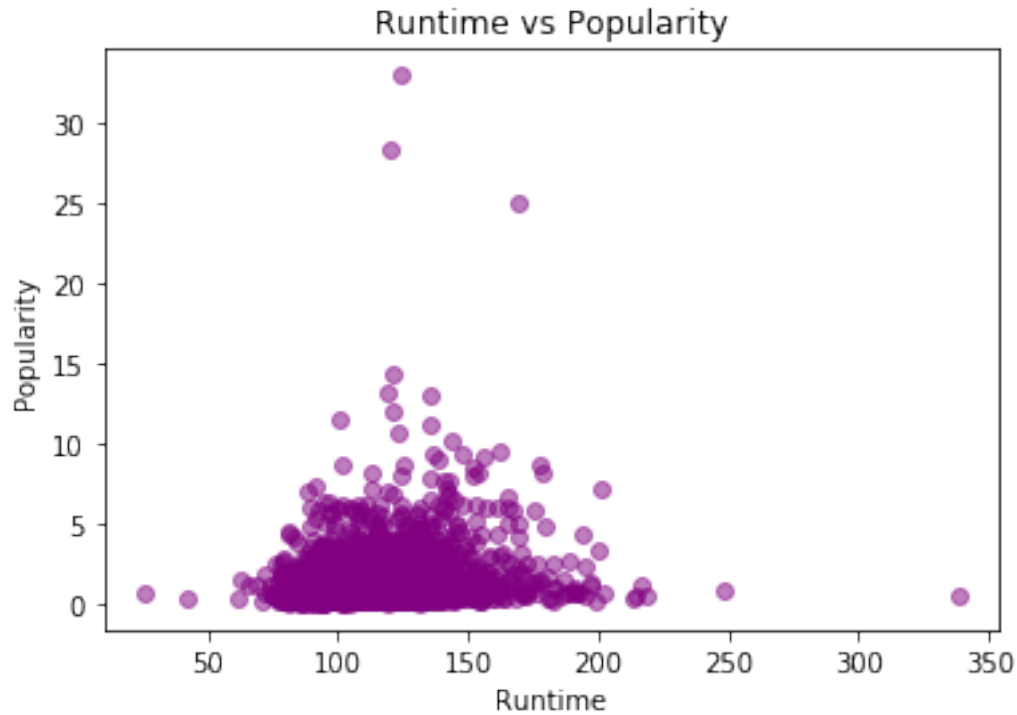
```
Out[52]: Text(0,0.5,'Number of Moives')
```



As we see from the bar chart the Drama and Comedy movies have the most popularity.

1.1.4 Research Question 2 (What the relationship between the runtime and the popularity of the movie?)

```
In [46]: # Create scatter and choose color
# Add title and axis names
plt.scatter(x=df['runtime'], y=df['popularity'],color='purple', alpha=0.5)
plt.title('Runtime vs Popularity')
plt.xlabel('Runtime')
plt.ylabel('Popularity')
plt.show()
```



We can see that the correlation coefficient is negative. That shows the relationship between runtime and popularity weak.

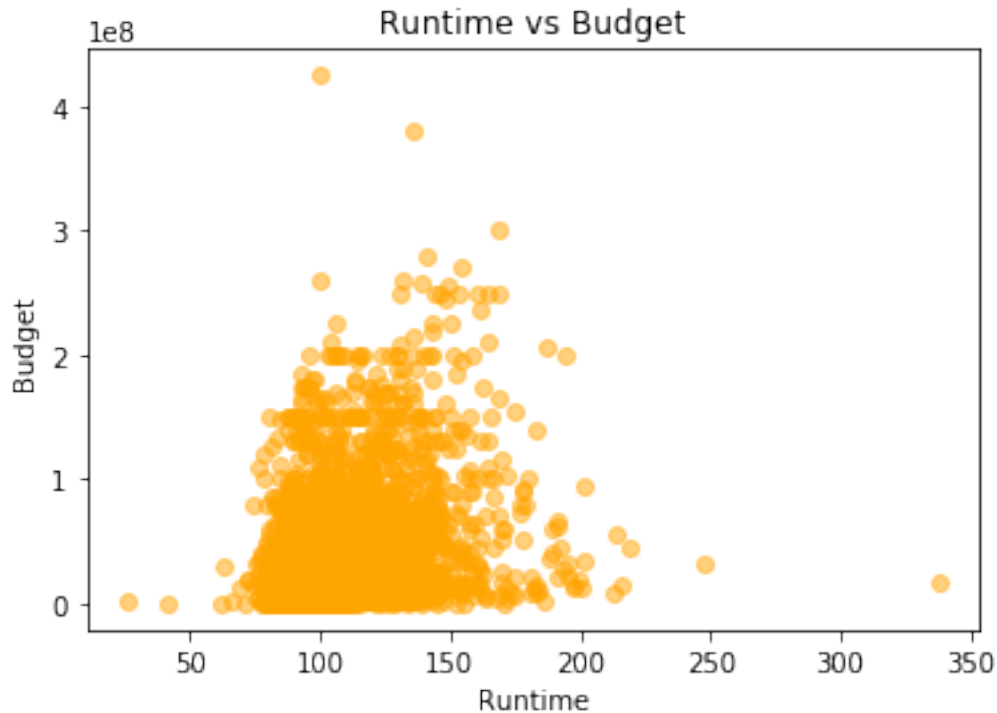
```
In [47]: df['runtime'].corr(df['popularity'])
```

```
Out[47]: 0.21342246569698353
```

This correlation is 0.213, a negative correlation between runtime and popularity.

Now, I want to know What the relationship between the runtime and the budget of the movie?

```
In [48]: # Create scatter and choose color
# Add title and axis names
plt.scatter(x=df['runtime'], y=df['budget'],color='orange', alpha=0.5)
plt.title('Runtime vs Budget')
plt.xlabel('Runtime')
plt.ylabel('Budget')
plt.show()
```



As we can see that the sign of the correlation coefficient is negative. That means the runtime decreases as budget increases.

```
In [49]: df['runtime'].corr(df['budget'])
```

```
Out[49]: 0.26079081142563248
```

Also, This correlation between runtime and budget is a negative correlation.

Conclusions

Finally: After I answered my questions. I came out with some facts about movies. After this analysis we can conclude the following: * The highest value of budget was 425000000 and in the profit was 2544505847. * The lowest value of the budget was zero and in profit was -413912431. * Drama movies and Comedy movies have the most popularity. * The correlation between runtime and budget is a negative correlation. * Also, the correlation between runtime and popularity is a negative correlation.

1.2 Limitations:

- In this dataset there something has hindered my analysis. At the budget and revenue column have some zero values in it. So that is erroneous and will adversely affect my overall analysis. I have removed the zero and null or missing values from the budget and revenue column, that makes the analysis more accurate.

1.3 Submitting your Project

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [50]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[50]: 0
```

```
In [ ]:
```